

Google File System

Google Disk Farm



Early days...

...today



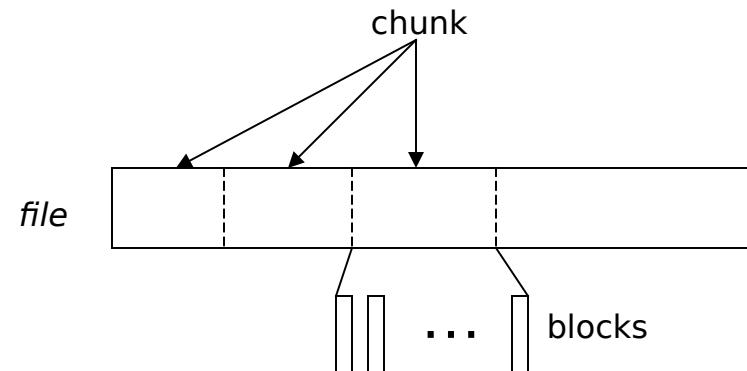
Design

■ Design factors


- Failures are common (built from inexpensive commodity components)
- Files
 - large (multi-GB)
 - mutation principally via appending new data
 - low-overhead atomicity essential
- Co-design applications and file system API
- Sustained bandwidth more critical than low latency

■ File structure

- Divided into 64 MB chunks
- Chunk identified by 64-bit handle
- Chunks replicated (default 3 replicas)
- Chunks divided into 64KB blocks
- Each block has a 32-bit checksum




Architecture




- Master
 - Manages namespace/metadata
 - Manages chunk creation, replication, placement
 - Performs snapshot operation to create duplicate of file or directory tree
 - Performs checkpointing and logging of changes to metadata
- Chunkservers
 - Stores chunk data and checksum for each block
 - On startup/failure recovery, reports chunks to master
 - Periodically reports sub-set of chunks to master (to detect no longer needed chunks)

Mutation operations

- Primary replica
 - Holds lease assigned by master (60 sec. default)
 - Assigns serial order for all mutation operations performed on replicas
- Write operation
 - 1-2: client obtains replica locations and identity of primary replica
 - 3: client pushes data to replicas (stored in LRU buffer by chunk servers holding replicas)
 - 4: client issues update request to primary
 - 5: primary forwards/perform write request
 - 6: primary receives replies from replica
 - 7: primary replies to client
- Record append operation
 - Performed atomically (one byte sequence)
 - At-least-once semantics
 - Append location chosen by GFS and returned to client
 - Extension to step 5:
 - If record fits in current chunk: write record and tell replicas the offset
 - If record exceeds chunk: pad the chunk, reply to client to use next chunk



Consistency Guarantees



- Write
 - Concurrent writes may be consistent but undefined
 - Write operations that are large or cross chunk boundaries are subdivided by client into individual writes
 - Concurrent writes may become interleaved
- Record append
 - Atomically, at-least-once semantics
 - Client retries failed operation
 - After successful retry, replicas are defined in region of append but may have intervening undefined regions
- Application safeguards
 - Use record append rather than write
 - Insert checksums in record headers to detect fragments
 - Insert sequence numbers to detect duplicates

| | Write | Record Append |
|----------------------|---|--|
| Serial success | <i>defined</i> | <i>defined</i> interspersed with <i>inconsistent</i> |
| Concurrent successes | <i>consistent</i> but <i>undefined</i> | |
| Failure | | <i>inconsistent</i> |

Metadata management

Logical structure →

| pathname | lock | chunk list |
|----------------|-------|---------------------|
| /home | read | Chunk4400488, ... |
| /save | | Chunk8ffe07783, ... |
| /home/user/foo | write | Chunk6254ee0, ... |
| /home/user | read | Chunk88f703, ... |

- Namespace
 - Logically a mapping from pathname to chunk list
 - Allows concurrent file creation in same directory
 - Read/write locks prevent conflicting operations
 - File deletion by renaming to a hidden name; removed during regular scan
- Operation log
 - Historical record of metadata changes
 - Kept on multiple remote machines
 - Checkpoint created when log exceeds threshold
 - When checkpointing, switch to new log and create checkpoint in separate thread
 - Recovery made from most recent checkpoint and subsequent log
- Snapshot
 - Revokes leases on chunks in file/directory
 - Log operation
 - Duplicate metadata (not the chunks!) for the source
 - On first client write to chunk:
 - Required for client to gain access to chunk
 - Reference count > 1 indicates a duplicated chunk
 - Create a new chunk and update chunk list for duplicate

Chunk/replica management

■ Placement

- On chunkservers with below-average disk space utilization
- Limit number of “recent” creations on a chunkserver (since access traffic will follow)
- Spread replicas across racks (for reliability)

■ Reclamation


- Chunk become garbage when file of which they are a part is deleted
- Lazy strategy (garbage college) is used since no attempt is made to reclaim chunks at time of deletion
- In periodic “HeartBeat” message chunkserver reports to the master a subset of its current chunks
- Master identifies which reported chunks are no longer accessible (i.e., are garbage)
- Chunkserver reclaims garbage chunks

■ Stale replica detection

- Master assigns a version number to each chunk/replica
- Version number incremented each time a lease is granted
- Replicas on failed chunkservers will not have the current version number
- Stale replicas removed as part of garbage collection

Performance

| Cluster | A | B |
|--------------------------|-------|--------|
| Chuckservers | 342 | 227 |
| Available disk space | 72 TB | 180 TB |
| Used disk space | 55 TB | 155 TB |
| Number of Files | 735 k | 737 k |
| Number of Dead files | 22 k | 232 k |
| Number of Chunks | 992 k | 1550 k |
| Metadata at chuckservers | 13 GB | 21 GB |
| Metadata at master | 48 MB | 60 MB |



| Cluster | A | B |
|----------------------------|-----------|-----------|
| Read rate (last minute) | 583 MB/s | 380 MB/s |
| Read rate (last hour) | 562 MB/s | 384 MB/s |
| Read rate (since restart) | 589 MB/s | 49 MB/s |
| Write rate (last minute) | 1 MB/s | 101 MB/s |
| Write rate (last hour) | 2 MB/s | 117 MB/s |
| Write rate (since restart) | 25 MB/s | 13 MB/s |
| Master ops (last minute) | 325 Ops/s | 533 Ops/s |
| Master ops (last hour) | 381 Ops/s | 518 Ops/s |
| Master ops (since restart) | 202 Ops/s | 347 Ops/s |