

MAP-REDUCE FRAMEWORK

-By Jagadish Rouniyar

ACTIVITIES

- ◎ 10 students are given the task of counting words in a book “Big Data”

How they will perform it ? Identify the step .
It denotes the map-reduce function

MAP - REDUCE

- sort/merge based distributed processing
 - Best for batch- oriented processing
 - Sort/merge is primitive
 - Operates at transfer rate (Process+data clusters)
 - Simple programming metaphor:
 - input | map | shuffle | reduce > output
 - cat * | grep | sort | uniq c > file
 - Pluggable user code runs in generic reusable framework
 - log processing,
 - web search indexing
 - SQL like queries in PIG
 - Distribution & reliability
 - Handled by framework - transparency

MR MODEL

- **Map()**

- Process a key/value pair to generate intermediate key/value pairs

- **Reduce()**

- Merge all intermediate values associated with the same key

- Users implement interface of two primary methods:

1. Map: $(key1, val1) \rightarrow (key2, val2)$
2. Reduce: $(key2, [val2]) \rightarrow [val3]$

- *Map - clause group-by (for Key) of an aggregate function of SQL*

- *Reduce - aggregate function (e.g., average) that is computed over all the rows with the same group-by attribute (key).*

- Application writer specifies
- A pair of functions called *Map* and *Reduce* and a set of input files and submits the job
- Workflow
- *Input* phase generates a number of *FileSplits* from input files (one per Map task)
- The *Map* phase executes a user function to transform input key-pairs into a new set of key-pairs
- The framework sorts & *Shuffles* the key-pairs to output nodes
- The *Reduce* phase combines all key-pairs with the same key into new keypairs
- The output phase writes the resulting pairs to files
- All phases are distributed with many tasks doing the work
- Framework handles scheduling of tasks on cluster
- Framework handles recovery when a node fails

Data distribution

- Input files are **split** into M pieces on distributed file system - 128 MB blocks
- Intermediate files created from *map* tasks are written to **local disk**
- Output files are written to **distributed file system**

Assigning tasks

- Many copies of user program are started
- Tries to utilize data **localization** by running *map* tasks on machines with data
- One instance becomes the Master
- Master finds idle machines and assigns them tasks

Master



Go and do
thy bidding

Workers



In progress



In progress



Idle



In progress



In progress




Idle

Execution

- Map workers read in contents of corresponding input partition
- Perform user-defined *map* computation to create intermediate <key,value> pairs
- Periodically buffered output pairs written to local disk

Reduce

- Reduce workers iterate over ordered intermediate data
 - Each unique key encountered - values are passed to user's reduce function
 - eg. <key, [value1, value2,..., valueN]>
- Output of user's *reduce* function is written to output file on global file system
- When all tasks have completed, master wakes up user program



Master Server distributes M map task to mappers and monitors their progress.



Map Worker reads the allocated data, saves the map results in local buffer.



Shuffle phase assigns reducers to these buffers, which are remotely read and processed by reducers.



Reducers o/p the result on stable storage.




Figure 1: Execution overview

DATA FLOW

- Input, final output are stored on a distributed file system
 - Scheduler tries to schedule map tasks “close” to physical storage location of input data
- Intermediate results are stored on local FS of map and reduce workers
- Output is often input to another map reduce task

COORDINATION

- Master data structures

- Task status: (idle, in-progress, completed)
- Idle tasks get scheduled as workers become available
- When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
- Master pushes this info to reducers

- Master pings workers periodically to detect failures

FAILURES

○ Map worker failure

- Map tasks completed or in-progress at worker are reset to idle
- Reduce workers are notified when task is rescheduled on another worker


○ Reduce worker failure

- Only in-progress tasks are reset to idle


○ Master failure

- MapReduce task is aborted and client is notified

EXECUTION



PARALLEL EXECUTION



HOW MANY MAP AND REDUCE JOBS?

- M map tasks, R reduce tasks
- Rule of thumb:
 - Make M and R much larger than the number of nodes in cluster
 - One DFS chunk per map is common
 - Improves dynamic load balancing and speeds recovery from worker failure
- Usually R is smaller than M, because output is spread across R files

COMBINERS

- Often a map task will produce many pairs of the form $(k, v_1), (k, v_2), \dots$ for the same key k
 - E.g., popular words in Word Count
- Can save network time by pre-aggregating at mapper
 - $\text{combine}(k_1, \text{list}(v_1)) \rightarrow v_2$
 - Usually same as reduce function
- Works only if reduce function is commutative and associative

PARTITION FUNCTION

- Inputs to map tasks are created by contiguous splits of input file
- For reduce, we need to ensure that records with the same intermediate key end up at the same worker
- System uses a default partition function e.g., $\text{hash}(\text{key}) \bmod R$
- Sometimes useful to override
 - E.g., $\text{hash}(\text{hostname(URL)}) \bmod R$ ensures URLs from a host end up in the same output file

EXECUTION SUMMARY


- How is this distributed?
 1. Partition input key/value pairs into chunks, run map() tasks in parallel
 2. After all map()s are complete, consolidate all emitted values for each unique emitted key
 3. Now partition space of output map keys, and run reduce() in parallel
- If map() or reduce() fails, reexecute!

Process	Time ----->
User Program	MapReduce() ... wait ...
Master	Assign tasks to worker machines...
Worker 1	Map 1 Map 3
Worker 2	Map 2
Worker 3	Read 1.1 Read 1.3 Read 1.2 Reduce 1
Worker 4	Read 2.1 Read 2.2 Read 2.3 Reduce 2

WORD COUNT EXAMPLE

Divide collection of document among the class.

Sum up the counts from all the documents to give final answer.



Each person gives count of individual word in a document. Repeats for assigned quota of documents.

(Done w/o communication)

```
map(String input_key, String input_value):  
    // input_key: document name  
    // input_value: document contents  
    for each word w in input_value:  
        EmitIntermediate(w, "1");  
reduce(String output_key, Iterator  
      intermediate_values):  
    // output_key: a word  
    // output_values: a list of counts  
    int result = 0;  
    for each v in intermediate_values:  
        result += ParseInt(v);  
    Emit(AsString(result));
```

○ Map()

- Input <filename, file text>
- Parses file and emits <word, count> pairs
 - eg. <"hello", 1>

○ Reduce()

- Sums all values for the same key and emits <word, TotalCount>
- eg. <"hello", (3 5 2 7)> => <"hello", 17>

- File

Hello World Bye World

Hello Hadoop GoodBye Hadoop

- Map

For the given sample input the first map emits:

< Hello, 1>

< World, 1>

< Bye, 1>

< World, 1>

- The second map emits:

< Hello, 1>

< Hadoop, 1>

< Goodbye, 1>

< Hadoop, 1>

The output of the first combine:


- < Bye, 1>
- < Hello, 1>
- < World, 2>






The output of the second combine:

- < Goodbye, 1>
- < Hadoop, 2>
- < Hello, 1>

Thus the output of the job (reduce) is:


- < Bye, 1>
- < Goodbye, 1>
- < Hadoop, 2>
- < Hello, 2>
- < World, 2>





MAP REDUCE ARCHITECTURE

SINGLE-NODE ARCHITECTURE



Machine Learning, Statistics

“Classical” Data Mining


COMMODITY CLUSTERS

- Web data sets can be very large
 - Tens to hundreds of terabytes
- Cannot mine on a single server (why?)
- Standard architecture emerging:
 - Cluster of commodity Linux nodes
 - Gigabit ethernet interconnect
- How to organize computations on this architecture?
 - Mask issues such as hardware failure

CLUSTER ARCHITECTURE

2-10 Gbps backbone between racks

1 Gbps between
any pair of nodes
in a rack




Each rack contains 16-64 nodes

STABLE STORAGE

- First order problem: if nodes can fail, how can we store data persistently?
- Answer: Distributed File System
 - Provides global file namespace
 - Google GFS; Hadoop HDFS; Kosmix KFS
- Typical usage pattern
 - Huge files (100s of GB to TB)
 - Data is rarely updated in place
 - Reads and appends are common

CONFIGURATION



HADOOP : MAP REDUCE FRAMEWORK

Hadoop Map-Reduce is a software framework for easily writing applications which process vast amounts of data (multi-terabyte data-sets) in-parallel on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.

FUNCTIONAL PROGRAMMING

- functional languages : Take example of Programming model from Lisp

MAP IN LISP (SCHEME)

⦿ (map *f* *list* [*list*₂ *list*₃ ...])

Unary operator

⦿ (map square '(1 2 3 4))

■ (1 4 9 16)

Binary operator

⦿ (reduce + '(1 4 9 16))

■ (+ 16 (+ 9 (+ 4 1)))

⦿ (reduce + (map square (map - l₁ l₂))))

FAULT TOLERANCE

- Workers are periodically pinged by master
 - No response = failed worker
- Master writes periodic checkpoints
- On errors, workers send “last gasp” UDP packet to master
 - Detect records that cause deterministic crashes and skips them

FAULT TOLERANCE

- Input file blocks stored on multiple machines
- When computation almost done, reschedule in-progress tasks

FAULT TOLERANCE / WORKERS

Handled via re-execution

- Detect failure via periodic heartbeats
- Re-execute completed + in-progress *map* tasks
 - Why????
- Re-execute in progress *reduce* tasks
- Task completion committed through master

MASTER FAILURE

- ◉ Could handle, ... ?
- ◉ But don't yet
 - (master failure unlikely)

REFINEMENT: REDUNDANT EXECUTION

Slow workers significantly delay completion time

- Other jobs consuming resources on machine
- Bad disks w/ soft errors transfer data slowly
- Weird things: processor caches disabled (!!)

Solution: Near end of phase, spawn backup tasks

- Whichever one finishes first "wins"

Dramatically shortens job completion time

REFINEMENT: LOCALITY OPTIMIZATION

- Master scheduling policy:

- Asks GFS for locations of replicas of input file blocks
- Map tasks typically split into 64MB (GFS block size)
- Map tasks scheduled so GFS input block replica are on same machine or same rack

- Effect

- Thousands of machines read input at local disk speed
 - Without this, rack switches limit read rate

REFINEMENT

SKIPPING BAD RECORDS

- ◉ Map/Reduce functions sometimes fail for particular inputs
 - Best solution is to debug & fix
 - Not always possible ~ third-party source libraries
 - On segmentation fault:
 - Send UDP packet to master from signal handler
 - Include sequence number of record being processed
 - If master sees two failures for same record:
 - Next worker is told to skip the record

OTHER REFINEMENTS

- ◉ Sorting guarantees
 - within each reduce partition
- ◉ Compression of intermediate data
- ◉ Combiner
 - Useful for saving network bandwidth
- ◉ Local execution for debugging/testing
- ◉ User-defined counters

PERFORMANCE

Tests run on cluster of 1800 machines:


- 4 GB of memory
 - Dual-processor 2 GHz Xeons with Hyperthreading
 - Dual 160 GB IDE disks
 - Gigabit Ethernet per machine
 - Bisection bandwidth approximately 100 Gbps

Two benchmarks:

MR_GrepScan 1010 100-byte records to extract records matching a rare pattern (92K matching records)

MR_SortSort 1010 100-byte records (modeled after TeraSort benchmark)

MR_GREP




Locality optimization helps:

- ⦿ 1800 machines read 1 TB at peak ~31 GB/s
- ⦿ W/out this, rack switches would limit to 10 GB/s



Startup overhead is significant for short jobs

MR_SORT

Normal



No backup tasks 200 processes killed



- Backup tasks reduce job completion time a lot!
- System deals well with failures

USAGE IN AUG 2004

Number of jobs	29,423
Average job completion time	634 secs
Machine days used	79,186 days
Input data read	3,288 TB
Intermediate data produced	758 TB
Output data written	193 TB
Average worker machines per job	157
Average worker deaths per job	1.2
Average map tasks per job	3,351
Average reduce tasks per job	55
Unique <i>map</i> implementations	395
Unique <i>reduce</i> implementations	269
Unique <i>map/reduce</i> combinations	426