



A Guide to the Java Search Engine

Lucene IN ACTION

Otis Gospodnetic
Erik Mecher
co-written by Jimmy Latimer



SEARCHING AND INDEXING BIG DATA

-By Jagadish Rouniyar

WHAT IS IT?

- Doug Cutting's grandmother's middle name
- A open source set of Java Classses
 - Search Engine/Document Classifier/Indexer
 - <http://lucene.sourceforge.net/talks/pisa/>
 - Developed by Doug Cutting 1996
 - Xerox/Apple/Excite/Nutch
 - Wrote several papers in IR


LUCENE

- Open source Java library for indexing and searching
 - Lets you add search to your application
 - High performance, scalable, full-text search library
 - Not a complete search system by itself
 - Written by Doug Cutting
- Used by LinkedIn, Twitter, ...
 - ...and many more (see
<http://wiki.apache.org/lucene-java/PoweredBy>)
- Ports/integrations to other languages
 - C/C++, C#, Ruby, Perl, Python, PHP, ...


RESOURCES

- Lucene: <http://lucene.apache.org/core/>
- Lucene in Action:
<http://www.manning.com/hatcher3/>
 - Code samples available for download
- Ant: <http://ant.apache.org/>
 - Java build system used by “Lucene in Action” code

LUCENE IN A SEARCH SYSTEM



LUCENE IMPLEMENTATION



LUCENE INDEX: CONCEPTS

The theory

Concepts

- Index: sequence of documents (a.k.a. Directory)
- Document: sequence of fields
- Field: named sequence of terms
- Term: a *text string* (e.g., a word)

Statistics

- Term frequencies and positions

CORE INDEXING CLASSES

- IndexWriter

- Central component that allows you to create a new index, open an existing one, and add, remove, or update documents in an index

- Directory


- Abstract class that represents the location of an index

- Analyzer

- Extracts tokens from a text stream

LUCENE IMPLEMENTATION

Application




lucene

LUCENE IMPLEMENTATION: INDEXING

Application


lucene



LUCENE IMPLEMENTATION: INDEXING

Application


Lucene



LUCENE IMPLEMENTATION: INDEXING


Application

lucene



LUCENE INDEXING


Application



Lucene

LUCENE INDEXING STEP 1 OF 5


Application



```
Document document = new Document();
document.add(new Field("title", title, Field.Store.YES,
                      Field.Index.TOKENIZED));
document.add(new Field("content", content));
document.add(new Field("id", id, Field.Store.YES,
                      Field.Index.NO));
try {
    indexWriter.addDocument(document);
} catch (IOException e) {
    e.printStackTrace();
}
```

LUCENE INDEXING STEP 2 OF 5


Application



```
Document document = new Document();
document.add(new Field("title", title, Field.Store.YES,
                      Field.Index.TOKENIZED));
document.add(new Field("content", content));
document.add(new Field("id", id, Field.Store.YES,
                      Field.Index.NO));
try {
    indexWriter.addDocument(document);
} catch (IOException e) {
    e.printStackTrace();
}
```

LUCENE INDEXING STEP 3 OF 5


Application



```
Document document = new Document();
document.add(new Field("title", title, Field.Store.YES,
                      Field.Index.TOKENIZED));
document.add(new Field("content", content));
document.add(new Field("id", id, Field.Store.YES,
                      Field.Index.NO));
try {
    indexWriter.addDocument(document);
} catch (IOException e) {
    e.printStackTrace();
}
```


LUCENE INDEXING STEP 4 OF 5

Application



LUCENE INDEXING STEP 5 OF 5

Application




Lucene

```
Document document = new Document();
document.add(new Field("title", title, Field.Store.YES,
                      Field.Index.TOKENIZED));
document.add(new Field("content", content));
document.add(new Field("id", id, Field.Store.YES,
                      Field.Index.NO));

try {
    indexWriter.addDocument(document);
} catch (IOException e) {
    e.printStackTrace();
}
```


LUCENE INDEXING

Application



LUCENE IMPLEMENTATION


Application



Lucene

LUCENE IMPLEMENTATION



Application



Lucene

SEARCHING:


Application



Lucene

SEARCHING: STEP 1 OF 6


Application



```
IndexSearcher indexSearcher = null;  
try{  
    indexSearcher = new IndexSearcher("/opt/lucene/index");  
}catch(IOException ioe){  
    ioe.printStackTrace();  
}
```


SEARCHING: STEP 2 OF 6

```
QueryParser queryParser = new QueryParser("content", analyzer);
```



SEARCHING: STEP 3 OF 6


Application



```
Query query = null;  
try {  
    query = queryParser.parse("Search string");  
} catch (ParseException e) {  
    e.printStackTrace();  
}
```

SEARCHING: STEP 4&5 OF 6


Application



```
if(null != query && null != indexSearcher) {  
    try {  
        Hits hits = indexSearcher.search(query);  
    } catch (IOException e) {  
        e.printStackTrace();  
    }  
}
```


SEARCHING: STEP 6 OF 6

Application



```
if(null != query && null != indexSearcher){  
    try {  
        Hits hits = indexSearcher.search(query);  
        for(int i = 0; i < hits.length(); i++){  
            System.out.print(hits.doc(i).get("id"));  
            System.out.println(hits.doc(i).get("title"));  
        }  
    } catch (IOException e) {  
        e.printStackTrace();  
    }  
}
```

SEARCHING



```

IndexSearcher indexSearcher = null;
try{
    indexSearcher = new IndexSearcher("/opt/lucene/index");
} catch(IOException ioe){
    ioe.printStackTrace();
}
QueryParser queryParser = new QueryParser("content", analyzer);
Query query = null;
try {
    query = queryParser.parse("Search string");
} catch (ParseException e) {
    e.printStackTrace();
}
if(null != query && null != indexSearcher){
    try {
        Hits hits = indexSearcher.search(query);
        for(int i = 0; i < hits.length(); i++){
            System.out.print(hits.doc(i).get("id"));
            System.out.println(hits.doc(i).get("title"));
        }
    } catch (IOException e) {
        e.printStackTrace();
    }
}
  
```

CREATING AN INDEXWRITER

```
import org.apache.lucene.index.IndexWriter;
import org.apache.lucene.store.Directory;
import org.apache.lucene.analysis.standard.StandardAnalyzer;
...
private IndexWriter writer;
...
public Indexer(String indexDir) throws IOException {
    Directory dir = FSDirectory.open(new File(indexDir));
    writer = new IndexWriter(
        dir,
        new
StandardAnalyzer(Version.LUCENE_30),
        true,
        IndexWriter.MaxFieldLength.UNLIMITED);
}
```

CORE INDEXING CLASSES (CONTD.)


- Document

- Represents a collection of named Fields. Text in these Fields are indexed.

- Field




- Note: Lucene Fields can represent both “fields” and “zones”

Lucene's data structures



Store

Documents:



A DOCUMENT CONTAINS FIELDS

```
import org.apache.lucene.document.Document;
import org.apache.lucene.document.Field;
...
protected Document getDocument(File f) throws Exception {
    Document doc = new Document();
    doc.add(new Field("contents", new FileReader(f)));
    doc.add(new Field("filename", f.getName(),
                      Field.Store.YES,
                      Field.Index.NOT_ANALYZED));
    doc.add(new Field("fullpath",
                      f.getCanonicalPath(),
                      Field.Store.YES,
                      Field.Index.NOT_ANALYZED));
    return doc;
}
```

INDEX A DOCUMENT WITH INDEXWRITER

```
private IndexWriter writer;  
...  
private void indexFile(File f) throws  
    Exception {  
    Document doc = getDocument(f);  
    writer.addDocument(doc);  
}
```

INDEXING A DIRECTORY

```
private IndexWriter writer;  
  
...  
public int index(String dataDir,  
                  FileFilter filter)  
    throws Exception {  
    File[] files = new File(dataDir).listFiles();  
    for (File f: files) {  
        if (... &&  
            (filter == null ||  
filter.accept(f))) {  
            indexFile(f);  
        }  
    }  
    return writer.numDocs();
```

CLOSING THE INDEXWRITER

```
private IndexWriter writer;  
...  
public void close() throws IOException  
{  
    writer.close();  
}
```

Indexing

○ **Attributes**

- Stored: original content retrievable
- Indexed: inverted, searchable
- Tokenized: analyzed, split into tokens

○ **Factory methods**

- Keyword: stored and indexed as single term
- Text: indexed, tokenized, and stored if String
- UnIndexed: stored
- UnStored: indexed, tokenized

○ **Terms are what matter for searching**

CORE SEARCHING CLASSES

- IndexSearcher
 - Central class that exposes several search methods on an index
- Query
 - Abstract query class. Concrete subclasses represent specific types of queries, e.g., matching terms in fields, boolean queries, phrase queries, ...
- QueryParser
 - Parses a textual representation of a query into a Query instance

CREATING AN INDEXSEARCHER

```
import org.apache.lucene.search.IndexSearcher;  
...  
public static void search(String indexDir,  
    String q)  
    throws IOException, ParseException {  
    Directory dir = FSDirectory.open(  
        new  
        File(indexDir));  
    IndexSearcher is = new IndexSearcher(dir);  
    ...  
}
```

QUERY AND QUERYPARSER

```
import org.apache.lucene.search.Query;
import org.apache.lucene.queryParser.QueryParser;
...
public static void search(String indexDir, String q)
    throws IOException, ParseException
...
QueryParser parser =
    new QueryParser(Version.LUCENE_30,
"contents",
    new StandardAnalyzer(
        Version.LUCENE_30));
```

CORE SEARCHING CLASSES (CONTD.)

- TopDocs
 - Contains references to the top documents returned by a search
- ScoreDoc
 - Represents a single search result

SEARCH() RETURNS TOPDOCS

```
import org.apache.lucene.search.TopDocs;  
...  
public static void search(String indexDir,  
                         String q)  
    throws IOException, ParseException  
...  
IndexSearcher is = ...;  
...  
Query query = ...;  
...  
TopDocs hits = is.search(query, 10);
```

TOPDOCS CONTAIN SCOREDOCS



```
import org.apache.lucene.search.ScoreDoc;  
...  
public static void search(String indexDir, String q)  
    throws IOException, ParseException  
...  
IndexSearcher is = ....;  
...  
TopDocs hits = ....;  
...  
for(ScoreDoc scoreDoc : hits.scoreDocs) {  
    Document doc = is.doc(scoreDoc.doc);  
    System.out.println(doc.get("fullpath"));  
}  
}
```

CLOSING INDEXSEARCHER

```
public static void search(String indexDir,  
                         String q)  
    throws IOException, ParseException  
{  
    ...  
    IndexSearcher is = ...;  
    ...  
    is.close();  
}
```

LUCENE CODE EXAMPLE: INDEXING

```
□ 01 Analyzer analyzer = new StandardAnalyzer();
02 IndexWriter iw = new IndexWriter("/tmp/testindex", analyzer, t
rue);
03
04 Document doc = new Document();
05 doc.add(new Field("body", "This is my TEST document",
Field.Store.YES, Field.Index.TOKENIZED));
06
07 iw.addDocument(doc);
08
09 iw.optimize();
10 iw.close();
```



StandardAnalyzer: **my, test, document**

LUCENE CODE EXAMPLE: SEARCHING

```
□ 01 Analyzer analyzer = new StandardAnalyzer();
02 IndexSearcher is = new IndexSearcher("/tmp/testindex");
03
04 QueryParser qp = new QueryParser("body", analyzer);
05 String userInput = "document AND test";
06 Query q = qp.parse(userInput);
07 Hits hits = is.search(q);
08 for (Iterator iter = hits.iterator(); iter.hasNext();) {
09     Hit hit = (Hit) iter.next();
10     System.out.println(hit.getScore() + " " + hit.get("body"));
11 }
12
13 is.close();
```

HOW LUCENE MODELS CONTENT

- A Document **is the atomic unit of indexing and searching**
 - A Document **contains Fields**
- **Fields have a name and a value**
 - You have to translate raw content into Fields
 - Examples: Title, author, date, abstract, body, URL, keywords, ...
 - Different documents can have different fields
 - Search a field using name:term, e.g., title:lucene

FIELDS

- Fields may
 - Be indexed or not
 - Indexed fields may or may not be analyzed (i.e., tokenized with an Analyzer)
 - Non-analyzed fields view the entire value as a single token (useful for URLs, paths, dates, social security numbers, ...)
 - Be stored or not
 - Useful for fields that you'd like to display to users
 - Optionally store term vectors
 - Like a positional index on the Field's terms
 - Useful for highlighting, finding similar documents, categorization

FIELD CONSTRUCTION

LOTS OF DIFFERENT CONSTRUCTORS

```
import org.apache.lucene.document.Field  
  
Field(String name,  
      String value,  
      Field.Store store, // store or  
not  
      Field.Index index, // index or  
not  
      Field.TermVector termVector);
```

**value can also be specified with a Reader, a
TokenStream, or a byte[]**

FIELD OPTIONS

○ Field.Store

- NO : Don't store the field value in the index
- YES : Store the field value in the index

○ Field.Index

- ANALYZED : Tokenize with an Analyzer
- NOT_ANALYZED : Do not tokenize
- NO : Do not index this field
- Couple of other advanced options

○ Field.TermVector

- NO : Don't store term vectors
- YES : Store term vectors
- Several other options to store positions and offsets

USING FIELD OPTIONS

Index	Store	TermVector	Example usage
NOT_ANALYZE D	YES	NO	Identifiers, telephone/SSNs, URLs, dates, ...
ANALYZED	YES	WITH_POSITIONS_OFFSETS	Title, abstract
ANALYZED	NO	WITH_POSITIONS_OFFSETS	Body
NO	YES	NO	Document type, DB keys (if not used for searching)
NOT_ANALYZE D	NO	NO	Hidden keywords

DOCUMENT

```
import org.apache.lucene.document.Field
```

- **Constructor:**

- Document();

- **Methods**

- void add(Fieldable field); // Field implements
// Fieldable
 - String get(String name); // Returns value of
Field with given
// name
 - Fieldable getFieldable(String name);
 - ... and many more

MULTI-VALUED FIELDS

- You can add multiple Fields with the same name
 - Lucene simply concatenates the different values for that named Field

```
Document doc = new Document();
doc.add(new Field("author",
                  "chris
manning",
                  Field.Store.YES,
                  Field.Index.ANALYZED));
doc.add(new Field("author",
                  "prabhakar
raghavan",
                  Field.Store.YES,
                  Field.Index.ANALYZED));
...
```

ANALYZERS

- Tokenizes the input text
- Common Analyzers
 - WhitespaceAnalyzer
Splits tokens on whitespace
 - SimpleAnalyzer
Splits tokens on non-letters, and then lowercases
 - StopAnalyzer
Same as SimpleAnalyzer, but also removes stop words
 - StandardAnalyzer
Most sophisticated analyzer that knows about certain token types, lowercases, removes stop words, ...

ANALYSIS EXAMPLES

- “The quick brown fox jumped over the lazy dog”
 - [The] [quick] [brown] [fox] [jumped] [over] [the] [lazy] [dog]
- WhitespaceAnalyzer
 - [the] [quick] [brown] [fox] [jumped] [over] [the] [lazy] [dog]
- SimpleAnalyzer
 - [the] [quick] [brown] [fox] [jumped] [over] [the] [lazy] [dog]
- StopAnalyzer
 - [quick] [brown] [fox] [jumped] [over] [lazy] [dog]
- StandardAnalyzer
 - [quick] [brown] [fox] [jumped] [over] [lazy] [dog]


MORE ANALYSIS EXAMPLES

- “XY&Z Corporation - xyz@example.com”
 - WhitespaceAnalyzer
 - [XY&Z] [Corporation] [-] [xyz@example.com]
 - SimpleAnalyzer
 - [xy] [z] [corporation] [xyz] [example] [com]
 - StopAnalyzer
 - [xy] [z] [corporation] [xyz] [example] [com]
 - StandardAnalyzer
 - [xy&z] [corporation] [xyz@example.com]

WHAT'S INSIDE AN ANALYZER?

- Analyzer's need to return a TokenStream

```
public TokenStream tokenStream(String  
fieldName,
```



TOKENIZERS AND TOKENFILTERS

Tokenizer

- WhitespaceTokenizer
- KeywordTokenizer
- LetterTokenizer
- StandardTokenizer
- ...

TokenFilter

- LowerCaseFilter
- StopFilter
- PorterStemFilter
- ASCIIFoldingFilter
- StandardFilter
- ...

INDEXWRITER CONSTRUCTION

```
// Deprecated  
IndexWriter(Directory d,  
           Analyzer a, //  
           default analyzer  
           IndexWriter.MaxFieldLength mfl);  
  
// Preferred  
IndexWriter(Directory d,  
           IndexWriterConfig c);
```

ADDING/DELETING DOCUMENTS TO/FROM AN INDEXWRITER

```
void addDocument(Document d);  
void addDocument(Document d, Analyzer a);
```

Important: Need to ensure that Analyzers used at indexing time are consistent with Analyzers used at searching time

```
// deletes docs containing term or matching  
// query. The term version is useful for  
// deleting one document.  
void deleteDocuments(Term term);  
void deleteDocuments(Query query);
```

INDEX FORMAT

- Each Lucene index consists of one or more segments
 - A segment is a standalone index for a subset of documents
 - All segments are searched
 - A segment is created whenever IndexWriter flushes adds/deletes
- Periodically, IndexWriter will merge a set of segments into a single segment
 - Policy specified by a MergePolicy
- You can explicitly invoke `optimize()` to merge segments

BASIC MERGE POLICY


- Segments are grouped into levels
- Segments within a group are roughly equal size (in log space)
- Once a level has enough segments, they are merged into a segment at the next level up

INDEXSEARCHER

◎ Constructor:

- IndexSearcher(Directory d);
 - deprecated

INDEXREADER



INDEXSEARCHER

◎ Constructor:

- `IndexSearcher(Directory d);`
 - **deprecated**
- `IndexSearcher(IndexReader r);`
 - **Construct an IndexReader with static method**
`IndexReader.open(dir)`

SEARCHING A CHANGING INDEX

```
Directory dir = FSDirectory.open(...);  
IndexReader reader = IndexReader.open(dir);  
IndexSearcher searcher = new IndexSearcher(reader);
```

Above reader does not reflect changes to the index unless you reopen it.

Reopening is more resource efficient than opening a new IndexReader.

```
IndexReader newReader = reader.reopen();  
If (reader != newReader) {  
    reader.close();  
    reader = newReader;  
    searcher = new IndexSearcher(reader);  
}
```

NEAR-REAL-TIME SEARCH

```
IndexWriter writer = ...;  
IndexReader reader = writer.getReader();  
IndexSearcher searcher = new IndexSearcher(reader);
```

Now let us say there's a change to the index using writer

```
// reopen() and getReader() force writer to flush  
IndexReader newReader = reader.reopen();  
if (reader != newReader) {  
    reader.close();  
    reader = newReader;  
    searcher = new IndexSearcher(reader);  
}
```

INDEXSEARCHER

◎ Methods

- TopDocs search(Query q, int n);
- Document doc(int docID);

QUERYPARSER

○ Constructor

- `QueryParser(Version matchVersion,
String
defaultField,
Analyzer
analyzer);`

○ Parsing methods

- `Query parse(String query) throws
ParseException;`
- ... and many more

QUERYPARSER SYNTAX EXAMPLES

Query expression	Document matches if...
java	Contains the term <i>java</i> in the default field
java junit java OR junit	Contains the term <i>java</i> or <i>junit</i> or both in the default field (<i>the default operator can be changed to AND</i>)
+java +junit java AND junit	Contains both <i>java</i> and <i>junit</i> in the default field
title:ant	Contains the term <i>ant</i> in the title field
title:extreme - subject:sports	Contains <i>extreme</i> in the title and not <i>sports</i> in subject
(agile OR extreme) AND java	Boolean expression matches
title:"junit in action"	Phrase matches in title
title:"junit action"~5	Proximity matches (within 5) in title
java*	Wildcard matches
java~	Fuzzy matches

CONSTRUCT QUERYS PROGRAMMATICALLY

- TermQuery
 - Constructed from a Term
- TermRangeQuery
- NumericRangeQuery
- PrefixQuery
- BooleanQuery
- PhraseQuery
- WildcardQuery
- FuzzyQuery
- MatchAllDocsQuery

TOPDOCS AND SCOREDOC

○ TopDocs methods

- Number of documents that matched the search
totalHits
- Array of ScoreDoc instances containing results
scoreDocs
- Returns best score of all matches
getMaxScore()

○ ScoreDoc methods

- Document id
doc
- Document score
score

SCORING

- Scoring function uses
 - Programmable boost values for certain fields in documents
 - Length normalization
 - Boosts for documents containing more of the query terms
- IndexSearcher provides an `explain()` method that explains the scoring of a document

SEARCH AND ANALYTICS (USING ELASTICSEARCH)

- Search - what's the big deal?
- Basic/Metadata retrieval
- “Find banks with more than (x) accounts”
- “Find banks *near my location*”



SEARCH CATEGORIES

- Basic/Metadata retrieval : data stores
- search engines:
 - Full-text search
 - Highlighting
 - Geolocation
 - Fuzzy search (“did-you-mean”)
 - Natural Language

ELASTICSEARCH

- **ElasticSearch** is a Distributed, RESTful, free/open source search server based on Apache Lucene. It is developed by Shay Banon^[1] and is released under the terms of the Apache License. ElasticSearch is developed in Java.
- Open-Source Search & Analytics engine
 - Structured & Unstructured Data
 - Real Time
 - Analytics capabilities (facets)
 - REST based
- Distributed
 - Designed for the Cloud
 - Designed for Big Data

ELASTICSEARCH

○ Advantages:

- ElasticSearch is distributed.
- ElasticSearch fully supports the near real-time search of Apache Lucene.
- Handling multitenancy is not a special configuration, whereas a more advanced setup is necessary with Solr.
- ElasticSearch introduces the concept of the Gateway, which makes full backups easier.

○ Disadvantages:

- Vendor dependent

USE CASE - TEXT SEARCH

The screenshot shows a GitHub search interface. The search bar at the top contains the query 'elastic'. Below the search bar, there's a sidebar with navigation links: 'PUBLIC', 'elastic', 'Open Source', 'http://elastic...', 'Clone', and dropdowns for 'branch: master' and 'Search GitHub for "elastic"'. The main content area displays search results for repositories and organizations:

- elasticsearch/elasticsearch-hadoop**: Read and write data to/from ElasticSearch with HDFS.
- elasticsearch/elasticsearch.github.com**
- elasticsearch/elasticsearch-hdfs**: Hadoop Plugin for ElasticSearch
- @elasticsearch**: Organization
- @elasticsearch-com**: Organization

Below the search results, there are two sections:

- elasticsearch /** 1000+ commits
 - Dates accessed from scripts should use UTC timezone ...
 - kimchy** authored 2 days ago latest commit b4d75a50bf
 - bin** 5 months ago Make elasticsearch.in.sh more configurable via env [philk]
 - config** 6 months ago Indexing Slow Log [kimchy]
 - lib** 2 years ago upgrade to sigar 1.6.4 [kimchy]
 - src** 2 days ago Dates accessed from scripts should use UTC timezone [kimchy]

USE CASE - GEOLOCATION

- Searches 50,000,000 venues every day using Elasticsearch

The screenshot displays the Foursquare website interface. At the top, there is a search bar with the text "restaurant" and a location input field set to "san francisco, ca". Below the search bar, a sidebar titled "Suggestions for restaurant" includes a "Show me places..." section with checkboxes for "I haven't been to", "My friends have been to", "I have been to before", and "With Foursquare specials". A promotional box encourages users to download the app with the text "Get the app to find great places on the go." and a "SEND" button. On the left side, there are two venue cards: "Bi-Rite Creamery" at 3692 18th St. (at Dolores St.) with a rating of 9.7, and "Humphry Slocombe" at 2790 Harrison St. (at 24th St.). Both cards show a thumbnail image of their respective products and a "Save" and "Like" button. The main right-hand portion of the screen is a map of the San Francisco Bay Area, with numerous blue location pins scattered across the city, each containing a number indicating the count of nearby venues. Labeled locations include Larkspur, Mill Valley, Tiburon, Richmond, El Cerrito, Berkeley, Emeryville, Piedmont, Oakland, Alameda, and Daly City.