

Searching and Indexing

Indexing

- *Indexing is the initial part of all search applications.*
- *Its goal is to process the original data into a highly efficient cross-reference lookup in order to facilitate rapid searching.*
- *The job is simple when the content is already textual in nature and its location is known.*

Indexing

- **Steps:**
- ***acquiring the content.***
 - *This process gathers and scopes the content that needs to be indexed.*
- ***build documents***
 - *The raw content that needs to be indexed has to be translated into the units (usually called documents) used by the search application.*
- ***document analysis***
 - The textual fields in a document cannot be indexed directly. Rather, the text has to be broken into a series of individual atomic elements called tokens.
 - This happens during the *document analysis step*. *Each token corresponds roughly to a word in the language*, and the analyzer determines how the textual fields in the document are divided into a series of tokens.
- ***index the document***
 - The final step is to *index the document*. *During the indexing step, the document is added to the index.*

Lucene

- Lucene is a free, open source project implemented in Java.
- licensed under Apache Software Foundation.
- Lucene itself is a single JAR (Java Archive) file, less than 1 MB in size, and with no dependencies, and integrates into the simplest Java stand-alone console program as well as the most sophisticated enterprise application.
- Rich and powerful full-text search library.
- Lucene to provide full-text indexing across both database objects and documents in various formats (Microsoft Office documents, PDF, HTML, text, and so on).
- supporting full-text search using Lucene requires two steps:
 - ***creating a lucence index***
 - *creating a lucence index* on the documents and/or database objects.
 - ***Parsing looking up***
 - *parsing* the user query and *looking up* the prebuilt index to answer the query.

Architecture




Figure 3.1: Typical components of search application architecture with Lucene components highlighted

Query: **not**

String comparison slow!

Solution: **Inverted index**

c:\docs\einstein.txt:




The important thing is not to
questioning.

c:\docs\shakespeare.txt:

To be or not to be.

Inverted index



Query: **not**

c:\docs\einstein.txt: 0
The important thing is not to stop questioning.

c:\docs\shakespeare.txt: 1
To be or not to be.

Inverted index

be	1
important	0
is	0
not	0 1
or	1
questioning	0
stop	0
to	0 1
the	0
thing	0

Document IDs

Query: "not to"

c:\docs\einstein.txt: 0
0 1 2 3 4 5
The important thing is not to
stop questioning.
6 7

c:\docs\shakespeare.txt: 1
0 1 2 3 4 5
To be or not to be.

Inverted index

be	1	1	5
important	0	1	
is	0	3	
not	0	4	1
or	1	2	
questioning	0	7	
stop	0	6	
to	0	5	1 0 4
the	0	0	
thing	0	2	


 Document IDs
 Positions

Query: "not to"

c:\docs\einstein.txt: 0
0 1 2 3 4 5
The important thing is not to
stop questioning.
6 7

c:\docs\shakespeare.txt: 1
0 1 2 3 4 5
To be or not to be.

Inverted index with Payloads



c:\docs\einstein.txt: 0
0 1 2 3 4 5
The important thing is not to
stop questioning.
6 7

c:\docs\shakespeare.txt: 1
0 1 2 3 4 5
To be or not to be.

- Document IDs
- Positions

Creating an Index (IndexWriter Class)

- *The first step in implementing full-text searching with Lucene is to build an index.*
- *To create an index, the first thing that need to do is to create an IndexWriter object.*
- *The IndexWriter object is used to create the index and to add new index entries (i.e., Documents) to this index. You can create an IndexWriter as follows*
- `IndexWriter indexWriter = new IndexWriter("index-directory", new StandardAnalyzer(), true);`

Parsing the Documents (Analyzer Class)

- The job of Analyzer is to "parse" each field of your data into indexable "tokens" or keywords.
- Several types of analyzers are provided out of the box. Table 1 shows some of the more interesting ones.
- **StandardAnalyzer**
 - A sophisticated general-purpose analyzer.
- **WhitespaceAnalyzer**
 - A very simple analyzer that just separates tokens using white space.
- **StopAnalyzer**
 - Removes common English words that are not usually useful for indexing.
- **SnowballAnalyzer**
 - An interesting experimental analyzer that works on word roots (a search on *rain* should also return entries with *raining*, *rained*, and so on).

Adding a Document/object to Index (Document Class)

- .To index an object, we use the Lucene Document class, to which we add the fields that you want indexed.
- Document doc = new Document();
- doc.add(new Field("description", hotel.getDescription(), Field.Store.YES, Field.Index.TOKENIZED));


Elasticsearch (NO-SQL)

What is ElasticSearch ?

- Open Source (No-sql DB)
- Distributed (cloud friendly)
- Highly-available
- Designed to speak JSON (JSON in, JSON out)

Highly available

- For each index you can specify:
- **Number of shards**
 - Each index has fixed number of shards
- **Number of replicas**
 - Each shard can have 0-many replicas, can be changed dynamically



A: { shards: 3, replicas: 2 }



B: { shards: 2, replicas: 3 }



C: { shards: 1, replicas: 0 }



Admin API

- Indices
 - Status
 - CRUD operation
 - Mapping, Open/Close, Update settings
 - Flush, Refresh, Snapshot, Optimize
- Cluster
 - Health
 - State
 - Node Info and stats
 - Shutdown

Rich query API

- There is rich Query DSL for search, includes:
- Queries
 - Boolean, Term, Filtered, MatchAll, ...
- Filters
 - And/Or/Not, Boolean, Missing, Exists, ...
- Sort
 - order:asc, order:desc
- Facets
 - Facets allows to provide aggregated data for the search request
 - Terms, Term, Terms_stat, Statistical,....

Scripting support

- There is a support for using scripting languages
- mvel (default)
- JS
- Groovy
- Python

ES_Query syntax

```
{  
  "fields": ["name", "Id"]  
}
```

```
{  
  "from": 0,  
  "size": 100,  
  "fields": ["name", "Id"]  
}
```

ES_Query syntax

```
{  
  "fields": ["name", "Id"]  
  "query": {  
  }  
}  
-----  
{"from":0, "size":100, "fields": ["name", "Id"],  
 "query": {  
   "filtered": {  
     "query": {  
       "match_all": {}  
     },  
     "filter": {  
       "terms": {  
         "name": ["ABC", "XYZ"]  
       }  
     }  
   }  
 }
```

Filter_syntax

```
{"filter": "terms": {  
    "name": ["ABC", "XYZ"]  
}
```

```
{  
"filter":  
{  
    "range":  
        {  
            "from": 0,  
            "to": 10  
        }  
    }  
}
```

Facets_Syntax

```
{  
“facets”:{  
    “facet_name”:{  
        “terms”:{  
            “name”:[“ABC”, “XYZ”]  
        }  
    }  
}
```

Facets_Syntax

```
{  
“facets”:{  
    “facet_name”:{  
        “term_stats”:{  
            “key_field”：“name”  
            “value_field”：“salary”  
        }  
    }  
}
```

Assignment

- *Indexing and Searching with LUCENE*