

# The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung  
SOSP '03

CSC 456 Operating Systems  
Seminar Presentation (11/11/2010)

Elif Eyigöz, Walter Lasecki

# Outline

---

- ▶ **Background**
- ▶ **Architecture**
- ▶ **Master**
- ▶ **Chunkserver**
- ▶ **Write/Read Algorithm**
- ▶ **Namespace management / Locking**
- ▶ **Reliability**
- ▶ **Replication in Chunkservers**
- ▶ **Replication in Master**
- ▶ **Conclusion**
- ▶ **Comparison with AFS and NFS**

# Background – Introduction

---

## Google

- ▶ Search engine: Huge workload
- ▶ Applications: Lots of data being processed
- ▶ Extreme scale: 100 TB of storage across 1000s of disks on over a thousand machines, accessed by hundreds of clients

## AFS, NFS versus GFS

Read/write, Cache, Replication, Consistency, Faults

# Background – Motivational Facts

---

- ▶ Extreme scale:
  - ▶ Thousands of commodity-class PC's
  - ▶ Multiple clusters distributed worldwide
  - ▶ Thousands of queries served per second
  - ▶ One query reads 100's of MB of data
  - ▶ One query consumes 10's of billions of CPU cycles
  - ▶ Google stores dozens of copies of the entire web
- 
- ▶ Goal: Large, distributed, highly fault tolerant file system

# Background – Design Observations

---


- ▶ Component failures are normal
- ▶ Fault tolerance and automatic recovery are needed
  
- ▶ Huge files are common (Multi GB)
- ▶ I/O and block size assumptions
  
- ▶ Record appends are more prevalent than random writes
- ▶ Appending is the focus of performance optimization and atomicity guarantees

# Outline

---

- ▶ Background
- ▶ **Architecture**
- ▶ Master
- ▶ Chunkserver
- ▶ Write/Read Algorithm
- ▶ Namespace management / Locking
- ▶ Reliability
- ▶ Replication in Chunkservers
- ▶ Replication in Master
- ▶ Conclusion
- ▶ Comparison with AFS and NFS

# Architecture



# The Role of the Master

---

Metadata Server : Maintain all file system metadata

- ▶ File namespace
- ▶ File to chunk mapping
- ▶ Chunk location information
- ▶ Keep all the metadata in the master's memory (FAST)

Location of chunks and their replicas:

- ▶ Master does not keep a persistent record (no disk I/O)
- ▶ Operations log for persistence and recovery
- ▶ Poll it from chunkservers (master as monitor, no extra cost)

# Master as metadata server

---

## Monitor

- ▶ *Heartbeat* messages to detect the state of each chunkserver
- ▶ Communicate with chunkservers periodically

## Centralized Controller

- ▶ System-wide activities
- ▶ Manage chunk replication in the whole system

# Master

---

- ▶ Single master
  - ▶ Simplify the design of the system
  - ▶ Control chunk placement using global knowledge
  - ▶ Potential bottleneck?
- 
- ▶ Avoiding bottleneck
    - ▶ Clients do no read/write data through master
    - ▶ Clients cache metadata (e.g., chunk handles, locations)
    - ▶ Client typically asks for multiple chunks in the same request
    - ▶ Master can also include the information for chunks immediately following those requested

# Caching

---

- ▶ Metadata

- ▶ Cached on client after being retrieved from the Master
- ▶ Only kept for a period of time to prevent stale data

- ▶ File caching

- ▶ Clients never cache file data
- ▶ Chunkservers never cache file data (Linux's buffer will for local files)
- ▶ File working sets too large
- ▶ Simplifies cache coherence

# Chunkserver

---

▶ Large chunk size

▶ 64 MB: much larger than typical file system blocks

▶ Advantages:

▶ reduces network overhead via reduced client-server interaction

▶ reduces network overhead

▶ reduces the size of metadata stored on master

▶ Disadvantages:


▶ internal fragmentation (solution: lazy space allocation)

▶ large chunks can become hot spots:

(solutions: higher replication factor, staggering application start time, client-to-client communication)


# Write/Read Algorithm

## ► Write Algorithm




# Write/Read Algorithm

## ► Write Algorithm




# Write/Read Algorithm

## ► Write Algorithm




# Write/Read Algorithm

## ► Write Algorithm




# Write/Read Algorithm

## ► Read Algorithm




# Write/Read Algorithm

## ► Read Algorithm




# Namespace Management

- ▶ Namespace
- ▶ How file system represent and manage the hierarchy of files and directories




# Namespace Management

- ▶ Traditional file systems
- ▶ Per-directory data structure that lists all files in the directory
- ▶ Hierarchy: based on *parent-child tree*



# Namespace Management

- ▶ Google File System
- ▶ No such per-directory data structure
- ▶ Lookup table mapping full pathname to metadata
- ▶ Hierarchy: based on full pathname



# Namespace Management

---

- ▶ What's the advantage of the namespace management in Google File System?
- ▶ Locking
- ▶ Transaction operation
- ▶ Deadlock prevention

# Locking – Transactional Operation

---

- ▶ Traditional file systems
- ▶ Locking on the certain data structure (e.g., mutex)
- ▶ Example: Create new file under a directory
- ▶ e.g., create /dir/file3

```
lock (dir->mutex)
create file3
add file3 to dir
unlock (dir->mutex)
```

- ▶ Only one creation in a certain directory at one time
  - Reason: lock the directory
  - Prevent from creating files with the same name simultaneously
  - Add new file into the directory

# Locking – Transactional Operation

## ▶ Google File System

### ▶ Locking on the pathname

### ▶ Example: Create new file under a directory

▶ e.g., create /dir/file3, /dir/file4/, .....

read_lock (/dir)	read_lock (/dir)	...
write_lock (/dir/file3)	write_lock (/dir/file4)	
create file3	create file4	
unlock (/dir/file3)	unlock (/dir/file4)	
unlock (/dir)	unlock (/dir)	

### ▶ Allow concurrent mutations in the same directory

□ Key: using read lock for dir. Why it can?


□ By locking pathname, it can lock the new file before it is created → Prevent from creating files with the same name simultaneously

▶ 24

□ No per-directory data structure → do not need to add new file into the directory

# Locking – Deadlock Prevention

- ▶ Traditional file systems (Linux)
- ▶ Hierarchy: based on *parent-child tree*
- ▶ Ordering for resources : parent, child
- ▶ Ordering for requests : lock parent first, then lock child
- ▶ Problem: e.g., rename operation



- ▶ First step of rename : lock dir1 and dir2
- ▶ But it cannot decide which one should be locked first
- ▶ Linux only allows one rename operation one time in one FS

# Locking – Deadlock Prevention

---

- ▶ Google file system
- ▶ Hierarchy: based on full pathname
- ▶ Based on consistent total order
- ▶ First ordered by level in the namespace tree
- ▶ Lexicographically within the same level
- ▶ Example: rename operation

```
Operation 1 : mv /root/dir1/file1 /root/dir2/file2  
Operation 2 : mv /root/dir1/file2 /root/dir2/file1
```

- ▶ First step of rename : lock dir1 and dir2
- ▶ Level order : dir1 and dir2 are in the same level, so then
- ▶ Lexographical order : Lock dir1 first

# Outline

---

- ▶ Background
- ▶ Architecture
- ▶ Master
- ▶ Chunkserver
- ▶ Write/Read Algorithm
- ▶ Namespace management / Locking
- ▶ **Reliability**
  - ▶ Replication in Chunkservers
  - ▶ Replication in Master
- ▶ Conclusion
- ▶ Comparison with AFS and NFS


# Replication in Chunkservers

---

- ▶ Replication
  - ▶ Create redundancy
  - ▶ Each chunk is replicated on multiple chunkservers
  - ▶ By default, store three replicas
- ▶ Replication of chunks
  - ▶ Each mutation is performed at all the chunk's replicas
  - ▶ Split the replication into two flows
    - ▶ Data flow : replicate the data
    - ▶ Control flow : replicate the write requests and their order

# Replication in Chunkservers


- ▶ Data flow (Step 3)
- ▶ Client pushes data to all replicas
- ▶ Data is pushed linearly along a chain of chunkservers in a pipelined fashion.
- ▶ e.g., Client → Sec. A → Primary → Sec. B
- ▶ Once a chunkserver receives some data, it starts forwarding immediately



# Replication in Chunkservers

## ▶ Control flow

- ▶ (Step 4) After data flow, client send write request to the primary
- ▶ The primary assigns consecutive serial number to all the write requests it received
- ▶ (Step 5) The primary forwards the write requests and the serial number order to all secondary replicas



# Replication in Master

---

- ▶ Replication of master state
- ▶ Log of all changes made to metadata
- ▶ Periodic checkpoints of the log
- ▶ Log and checkpoints are replicated on multiple machines
- ▶ If master fails, another replica would be activated as the new master
- ▶ “shadow” masters provide read-only service
- ▶ Not mirrors, may lag the master slightly
- ▶ Enhance read throughput and availability for some files

# Outline

---

- ▶ Background
- ▶ Architecture
- ▶ Master
- ▶ Chunkserver
- ▶ Write/Read Algorithm
- ▶ Namespace management / Locking
- ▶ Reliability
- ▶ Replication in Chunkservers
- ▶ Replication in Master
- ▶ **Conclusion**
- ▶ Comparison with AFS and NFS

# Conclusion

---

- ▶ AFS, NFS and GFS
  - ▶ Read/write, Cache, Replication, Consistency, Faults
- ▶ The design is driven by Google's application workloads and technological environment
- ▶ Advantages
  - ▶ Know how to improve performance
  - ▶ Design the file system only for Google
- ▶ Disadvantages
  - ▶ May not be adopted in other system
  - ▶ Single master may be still a potential bottleneck

# References

---

- ▶ [1] The Google File System, SOSP '03
- ▶ [2] Presentation slide of GFS in SOSP '03
- ▶ [3] “Operating System Concepts”, 8th edition, 2009
- ▶ [4] Linux 2.6.20

---

Thank you!

---

